

[8/6, 09:00] +91 96866 77992: Q1) What is AWS?

AWS stands for Amazon Web Services. AWS is a platform that provides on-demand resources for hosting web services, storage, networking, databases and other resources over the internet with a pay-as-you-go pricing.

Q2) What are the components of AWS?

EC2 – Elastic Compute Cloud, S3 – Simple Storage Service, Route53, EBS – Elastic Block Store, Cloudwatch, Key-Pairs are few of the components of AWS.

Q3) What are key-pairs?

Key-pairs are secure login information for your instances/virtual machines. To connect to the instances we use key-pairs that contain a public-key and private-key.

Q4) What is S3?

S3 stands for Simple Storage Service. It is a storage service that provides an interface that you can use to store any amount of data, at any time, from anywhere in the world. With S3 you pay only for what you use and the payment model is pay-as-you-go.

Q5) What are the pricing models for EC2 instances?

The different pricing model for EC2 instances are as below,

On-demand Reserved Spot Scheduled Dedicated
Q6) What are the types of volumes for EC2 instances?
There are two types of volumes, Instance store volumes EBS – Elastic Block Stores
Q7) What are EBS volumes?

EBS stands for Elastic Block Stores. They are persistent volumes that you can attach to the instances. With EBS volumes, your data will be preserved even when you stop your instances, unlike your instance store volumes where the data is deleted when you stop the instances.

Q8) What are the types of volumes in EBS?

Following are the types of volumes in EBS,

General purpose Provisioned IOPS Magnetic Cold HDD Throughput optimized Q9) What are the different types of instances?

Following are the types of instances,

General purpose Computer Optimized Storage Optimized Memory Optimized Accelerated Computing Q10) What is an auto-scaling and what are the components?

Auto scaling allows you to automatically scale-up and scale-down the number of instances depending on the CPU utilization or memory utilization. There are 2 components in Auto scaling, they are Auto-scaling groups and Launch Configuration.

Q11) What are reserved instances?

Reserved instances are the instance that you can reserve a fixed capacity of EC2 instances. In reserved instances you will have to get into a contract of 1 year or 3 years.

Q12) What is an AMI?

AMI stands for Amazon Machine Image. AMI is a template that contains the software configurations, launch permission and a block device mapping that specifies the volume to attach to the instance when it is launched.

Q13) What is an EIP?

EIP stands for Elastic IP address. It is designed for dynamic cloud computing. When you want to have a static IP address for your instances when you stop and restart your instances, you will be using EIP address.

Q14) What is Cloudwatch?

Cloudwatch is a monitoring tool that you can use to monitor your various AWS resources. Like health check, network, Application, etc.

Q15) What are the types in cloudwatch?

There are 2 types in cloudwatch. Basic monitoring and detailed monitoring. Basic monitoring is free and detailed monitoring is chargeable.

Q16) What are the cloudwatch metrics that are available for EC2 instances?

Diskreads, Diskwrites, CPU utilization, networkpacketsIn, networkpacketsOut, networkIn, networkOut, CPUCreditUsage, CPUCreditBalance.

Q17) What is the minimum and maximum size of individual objects that you can store in S3

The minimum size of individual objects that you can store in S3 is 0 bytes and the maximum bytes that you can store for individual objects is 5TB.

Q18) What are the different storage classes in S3?

Following are the types of storage classes in S3,

Standard frequently accessed
Standard infrequently accessed
One-zone infrequently accessed
Glacier
RRS – reduced redundancy storage

Q19) What is the default storage class in S3?

The default storage class in S3 is Standard frequently accessed.

Q20) What is glacier?

Glacier is the back up or archival tool that you use to back up your data in S3.

Q21) How can you secure the access to your S3 bucket?

There are two ways that you can control the access to your S3 buckets,

ACL – Access Control List
Bucket policies

Q22) How can you encrypt data in S3?

You can encrypt the data by using the below methods,

Server Side Encryption – S3 (AES 256 encryption)
Server Side Encryption – KMS (Key management Service)
Server Side Encryption – C (Client Side)

Q23) What are the parameters for S3 pricing?

The pricing model for S3 is as below,

Storage used
Number of requests you make
Storage management
Data transfer
Transfer acceleration

Q24) What is the pre-requisite to work with Cross region replication in S3?

You need to enable versioning on both source bucket and destination to work with cross region replication. Also both the source and destination bucket should be in different region.

Q25) What are roles?

Roles are used to provide permissions to entities that you trust within your AWS account. Roles are users in another account. Roles are similar to users but with roles you do not need to create any username and password to work with the resources.

Q26) What are policies and what are the types of policies?

Policies are permissions that you can attach to the users that you create. These policies will contain that access that you have provided to the users that you have created. There are 2 types of policies.

Managed policies
Inline policies

Q27) What is cloudfront?

Cloudfront is an AWS web service that provided businesses and application developers an easy and efficient way to distribute their content with low latency and high data transfer speeds. Cloudfront is content delivery network of AWS.

Q28) What are edge locations?

Edge location is the place where the contents will be cached. When a user tries to access some

content, the content will be searched in the edge location. If it is not available then the content will be made available from the origin location and a copy will be stored in the edge location.

Q29) What is the maximum individual archive that you can store in glacier?

You can store a maximum individual archive of upto 40 TB.

Q30) What is VPC?

VPC stands for Virtual Private Cloud. VPC allows you to easily customize your networking configuration. VPC is a network that is logically isolated from other network in the cloud. It allows you to have your own IP address range, subnets, internet gateways, NAT gateways and security groups.

Q31) What is VPC peering connection?

VPC peering connection allows you to connect 1 VPC with another VPC. Instances in these VPC behave as if they are in the same network.

Q32) What are NAT gateways?

NAT stands for Network Address Translation. NAT gateways enables instances in a private subnet to connect to the internet but prevent the internet from initiating a connection with those instances.

Q33) How can you control the security to your VPC?

You can use security groups and NACL (Network Access Control List) to control the security to your

VPC.

Q34) What are the different types of storage gateway?

Following are the types of storage gateway.

File gateway Volume gateway Tape gateway Q35) What is a snowball?

Snowball is a data transport solution that used source appliances to transfer large amounts of data into and out of AWS. Using snowball, you can move huge amount of data from one place to another which reduces your network costs, long transfer times and also provides better security.

Q36) What are the database types in RDS?

Following are the types of databases in RDS,

Aurora Oracle MySQL server Postgresql MariaDB SQL server Q37) What is a redshift?

Amazon redshift is a data warehouse product. It is a fast and powerful, fully managed, petabyte scale data warehouse service in the cloud.

Q38) What is SNS?

SNS stands for Simple Notification Service. SNS is a web service that makes it easy to notifications from the cloud. You can set up SNS to receive email notification or message notification.

Q39) What are the types of routing policies in Route 53?

Following are the types of routing policies in Route 53,

Simple routing Latency routing Failover routing Geolocation routing Weighted routing Multivalue
answer Q40) What is the maximum size of messages in SQS?

The maximum size of messages in SQS is 256 KB.

Q41) What are the types of queues in SQS?

There are 2 types of queues in SQS.

Standard queue FIFO (First In First Out) Q42) What is Multi-AZ RDS?

Multi-AZ (Availability Zone) RDS allows you to have a replica of your production database in another availability zone. Multi-AZ (Availability Zone) database is used for disaster recovery. You will have an exact copy of your database. So when your primary database goes down, your application will automatically failover to the standby database.

Q43) What are the types of backups in RDS database?

There are 2 types of backups in RDS database.

Automated backups Manual backups which are known as snapshots. Q44) What is the difference between security groups and network access control list? Security Groups Network access control list Can control the access at the instance level Can control access at the subnet

levelCan add rules for "allow" onlyCan add rules for both "allow" and "deny"Evaluates all rules before allowing the trafficRules are processed in order number when allowing traffic.Can assign unlimited number of security groupsCan assign upto 5 security groups.Statefull filteringStateless filtering Q45) What are the types of load balancers in EC2?

There are 3 types of load balancers,

Application load balancerNetwork load balancerClassic load balancer Q46) What is and ELB?

ELB stands for Elastic Load balancing. ELB automatically distributes the incoming application traffic or network traffic across multiple targets like EC2, containers, IP addresses.

Q47) What are the two types of access that you can provide when you are creating users?

Following are the two types of access that you can create.

Programmatic accessConsole access Q48) What are the benefits of auto scaling?

Following are the benefits of auto scaling

Better fault toleranceBetter availabilityBetter cost management Q49) What are security groups?

Security groups acts as a firewall that contains the traffic for one or more instances. You can associate one or more security groups to your instances when you launch them. You can add rules to each security group that allow traffic to and from its associated instances. You can modify the rules of a security group at any time, the new rules are automatically and immediately applied to all the instances that are associated with the security group

Q50) What are shared AMI's?

Shared AMI's are the AMI that are created by other developed and made available for other developed to use.

Q51) What is the difference between the classic load balancer and application load balancer?

Answer: Dynamic port mapping, multiple port multiple listeners is used in Application Load Balancer, One port one listener is achieved via Classic Load Balancer

Q52) By default how many Ip address does aws reserve in a subnet?

Answer: 5

Q53) What is meant by subnet?

Answer: A large section of IP Address divided in to chunks are known as subnets

Q54) How can you convert a public subnet to private subnet?

Answer: Remove IGW & add NAT Gateway, Associate subnet in Private route table

Q55) Is it possible to reduce a ebs volume?

Answer: no it's not possible, we can increase it but not reduce them

Q56) What is the use of elastic ip are they charged by AWS?

Answer: These are ipv4 address which are used to connect the instance from internet, they are charged if the instances are not attached to it

Q57) One of my s3 bucket is deleted but i need to restore is there any possible way?

Answer: If versioning is enabled we can easily restore them

Q58) When I try to launch an ec2 instance i am getting Service limit exceed, how to fix the issue?

Answer: By default AWS offer service limit of 20 running instances per region, to fix the issue we need to contact AWS support to increase the limit based on the requirement

Q59) I need to modify the ebs volumes in Linux and windows is it possible

Answer: yes its possible from console use modify volumes in section give the size u need then for windows go to disk management for Linux mount it to achieve the modification

Q60) Is it possible to stop a RDS instance, how can I do that?

Answer: Yes it's possible to stop rds. Instance which are non-production and non multi AZ's

Q61) What is meant by parameter groups in rds. And what is the use of it?

Answer: Since RDS is a managed service AWS offers a wide set of parameter in RDS as parameter group which is modified as per requirement

Q62) What is the use of tags and how they are useful?

Answer: Tags are used for identification and grouping AWS Resources

Q63) I am viewing an AWS Console but unable to launch the instance, I receive an IAM Error how can I rectify it?

Answer: As AWS user I don't have access to use it, I need to have permissions to use it further

Q64) I don't want my AWS Account id to be exposed to users how can I avoid it?

Answer: In IAM console there is option as sign in url where I can rename my own account name with AWS account

Q65) By default how many Elastic Ip address does AWS Offer?

Answer: 5 elastic ip per region

Q66) You are enabled sticky session with ELB. What does it do with your instance?

Answer: Binds the user session with a specific instance

Q67) Which type of load balancer makes routing decisions at either the transport layer or the Application layer and supports either EC2 or VPC.

Answer: Classic Load Balancer

Q68) Which is virtual network interface that you can attach to an instance in a VPC?

Answer: Elastic Network Interface

Q69) You have launched a Linux instance in AWS EC2. While configuring security group, you have selected SSH, HTTP, HTTPS protocol. Why do we need to select SSH?

Answer: To verify that there is a rule that allows traffic from EC2 Instance to your computer

Q70) You have chosen a windows instance with Classic and you want to make some change to the security group. How will these changes be effective?

Answer: Changes are automatically applied to windows instances

Q71) Load Balancer and DNS service comes under which type of cloud service?

Answer: IaaS-Storage

Q72) You have an EC2 instance that has an unencrypted volume. You want to create another Encrypted volume from this unencrypted volume. Which of the following steps can achieve this?

Answer: Create a snapshot of the unencrypted volume (applying encryption parameters), copy the. Snapshot and create a volume from the copied snapshot

Q73) Where does the user specify the maximum number of instances with the auto scaling Commands?

Answer: Auto scaling Launch Config

Q74) Which are the types of AMI provided by AWS?

Answer: Instance Store backed, EBS Backed

Q75) After configuring ELB, you need to ensure that the user requests are always attached to a Single instance. What setting can you use?

Answer: Sticky session

Q76) When do I prefer to Provisioned IOPS over the Standard RDS storage?

Ans: If you have do batch-oriented is workloads.

Q77) If I am running on my DB Instance a Multi-AZ deployments, can I use to the stand by the DB Instance for read or write a operation along with to primary DB instance?

Ans: Primary db instance does not working.

Q78) Which the AWS services will you use to the collect and the process e-commerce data for the near by real-time analysis?

Ans: Good of Amazon DynamoDB.

Q79) A company is deploying the new two-tier an web application in AWS. The company has to limited on staff and the requires high availability, and the application requires to complex

queries and table joins. Which configuration provides to the solution for company's requirements?

Ans: An web application provide on Amazon DynamoDB solution.

Q80) Which the statement use to cases are suitable for Amazon DynamoDB?

Ans: The storing metadata for the Amazon S3 objects& The Running of relational joins and complex an updates.

Q81) Your application has to the retrieve on data from your user's mobile take every 5 minutes and then data is stored in the DynamoDB, later every day at the particular time the data is an extracted into S3 on a per user basis and then your application is later on used to visualize the data to user. You are the asked to the optimize the architecture of the backend system can to lower cost, what would you recommend do?

Ans: Introduce Amazon ElastiCache to the cache reads from the Amazon DynamoDB table and to reduce the provisioned read throughput.

Q82) You are running to website on EC2 instances can deployed across multiple Availability Zones with an Multi-AZ RDS MySQL Extra Large DB Instance etc. Then site performs a high number of the small reads and the write per second and the relies on the eventual consistency model. After the comprehensive tests you discover to that there is read contention on RDS MySQL. Which is the best approaches to the meet these requirements?

Ans:The Deploy Elasti Cache in-memory cache is running in each availability zone and Then Increase the RDS MySQL Instance size and the Implement provisioned IOPS.

Q83) An startup is running to a pilot deployment of around 100 sensors to the measure street noise and The air quality is urban areas for the 3 months. It was noted that every month to around the 4GB of sensor data are generated. The company uses to a load balanced take auto

scaled layer of the EC2 instances and a RDS database with a 500 GB standard storage. The pilot was success and now they want to the deploy take atleast 100K sensors.let which to need the supported by backend. You need to the stored data for at least 2 years to an analyze it. Which setup of following would you be prefer?

Ans: The Replace the RDS instance with an 6 node Redshift cluster with take 96TB of storage.

Q84) Let to Suppose you have an application where do you have to render images and also do some of general computing. which service will be best fit your need?

Ans: Used on Application Load Balancer.

Q85) How will change the instance give type for the instances, which are the running in your applications tier and Then using Auto Scaling. Where will you change it from areas?

Ans: Changed to Auto Scaling launch configuration areas.

Q86) You have an content management system running on the Amazon EC2 instance that is the approaching 100% CPU of utilization. Which option will be reduce load on the Amazon EC2 instance?

Ans: Let Create a load balancer, and Give register the Amazon EC2 instance with it.

Q87) What does the Connection of draining do?

Ans: The re-routes traffic from the instances which are to be updated (or) failed an health to check.

Q88) When the instance is an unhealthy, it is do terminated and replaced with an new ones,

which of the services does that?

Ans: The service makes a fault tolerance.

Q89) What are the life cycle hooks used for the AutoScaling?

Ans: They are used to put an additional taken wait time to the scale in or scale out events.

Q90) An user has to setup an Auto Scaling group. Due to some issue the group has failed to launch a single instance for more than 24 hours. What will happen to the Auto Scaling in the condition?

Ans: The auto Scaling will be suspended to the scaling process.

Q91) You have an EC2 Security Group with several running EC2 instances. You changed the Security Group rules to allow the inbound traffic on a new port and protocol, and then launched several new instances in the same Security Group. Such the new rules apply?

Ans: The immediately to all the instances in security groups.

[8/6, 09:00] +91 96866 77992: What is Amazon Cloudwatch?

Answer :

CloudWatch is a monitoring service for AWS cloud resources and the applications you run on AWS. You can use CloudWatch to collect and track metrics, collect and monitor log files, and set alarms. CloudWatch can monitor resources such as EC2 instances, DynamoDB tables, and RDS DB instances.

Question 2. Which Operating Systems Does Cloudwatch Support?

Answer :

CloudWatch receives and provides metrics for all EC2 instances and should work with any operating system currently supported by the EC2 service.

Question 3. What Access Management Policies Can I Implement For Cloudwatch?

Answer :

CloudWatch integrates with AWS IAM so that you can specify which CloudWatch actions a user in your AWS Account can perform. You can't use IAM to control access to CloudWatch data for specific resources. For example, you can't give a user access to CloudWatch data for only a specific set of instances or a specific LoadBalancer. Permissions granted using IAM cover all the cloud resources you use with CloudWatch. In addition, you can't use IAM roles with the Amazon CloudWatch command line tools.

Question 4. What Is Amazon Cloudwatch Logs?

Answer :

Amazon CloudWatch Logs lets you monitor and troubleshoot your systems and applications using your existing system, application and custom log files. With CloudWatch Logs, you can monitor your logs, in near real time, for specific phrases, values or patterns. For example, you could set an alarm on the number of errors that occur in your system logs or view graphs of latency of web requests from your application logs. You can then view the original log data to see the source of the problem. Log data can be stored and accessed indefinitely in highly durable, low-cost storage so you don't have to worry about filling up hard drives.

Question 5. What Kinds Of Things Can I Do With Cloudwatch Logs?

Answer :

CloudWatch Logs is capable of monitoring and storing your logs to help you better understand and operate your systems and applications.

You can use CloudWatch logs in a number of ways:

Real Time application and system monitoring, Long term log retention.

[8/6, 09:00] +91 96866 77992: uestion 6. What Platforms Does The Cloudwatch Logs Agent Support?

Answer :

The CloudWatch Log Agent is supported on Amazon Linux, Ubuntu, CentOS, Red Hat Enterprise Linux, and Windows. This agent will support the ability to monitor individual log files on the host.

Question 7. Does The Cloudwatch Logs Agent Support Iam Roles?

Answer :

Yes. The CloudWatch Logs Agent is integrated with IAM and includes support for both access keys and IAM roles.

Question 8. Does The Amazon Cloudwatch Monitoring Charge Change Depending On Which Ec2 Instance I Monitor?

Answer :

No, CloudWatch monitoring charge does not vary by EC2 instance type.

Question 9. What Can I Measure With Amazon Cloudwatch Metrics?

Answer :

CloudWatch allows you to monitor AWS cloud resources and the applications you run on AWS. Metrics are provided automatically for a number of AWS products and services, including EC2 instances, EBS volumes, ELBs, Autoscaling groups, EMR job flows, RDS DB instances, DynamoDB tables, ElastiCache clusters, RedShift clusters, OpsWorks stacks, Route 53 health checks, SNS topics, SQS queues, SWF workflows, and Storage Gateways. You can also monitor custom metrics generated by your own applications and services.

Question 10. What Is The Retention Period Of All Metrics?

Answer :

CloudWatch retains metric data as follows:

Data points with a period of less than 60 seconds are available for 3 hours. These data points are high-resolution custom metrics, Data points with a period of 60 seconds (1 minute) are available for 15 days, Data points with a period of 300 seconds (5 minute) are available for 63 days, 4) Data points with a metric of 3600 seconds (1 hr) are available for 455 days (15 months). Data points that are initially published with a shorter period are aggregated together for long-term storage.

[8/6, 09:01] +91 96866 77992: Will I Lose The Metrics Data If I Disable Monitoring For An Ec2 Instance?

Answer :

No. You can always retrieve metrics data for any EC2 instance based on the retention schedules. However, the CloudWatch console limits the search of metrics to 2 weeks after a metric is last ingested to ensure that the most up to date instances are shown in your namespace.

Question 12. Can I Access The Metrics Data For A Terminated Ec2 Instance Or A Deleted Elb?

Answer :

Yes, CloudWatch stores metrics for terminated EC2 instances or deleted ELBs for 15 months.

Question 13. What Metrics Are Available At High Resolution?

Answer :

Currently, only customer metrics that you publish to CloudWatch are available at high resolution.

Question 14. Are High-resolution Custom Metrics Priced Differently Than Regular Custom Metrics?

Answer :

No, high-resolution custom metrics are priced in the same manner as standard 1-min custom metrics.

Question 15. When Would I Use A Custom Metric Over Having My Program Emit A Log To Cloudwatch Logs?

Answer :

You can monitor your own data using custom metrics, CloudWatch logs, or both. You may want to use custom metrics if your data is not already produced in log format, for example OS process or performance measurements. Or, you may want to write your own application or script, or one provided by an AWS partner. If you want to store and save individual measurements along with additional detail, you may want to use CloudWatch Logs.

Question 16. What Statistics Can I View And Graph In Cloudwatch?

Answer :

You can retrieve, graph, and set alarms on the following statistical metrics: average, sum, min, max, and sample count.

[8/6, 09:01] +91 96866 77992: 17. What Log Monitoring Does Cloudwatch Provide?

Answer :

CloudWatch Logs lets you monitor and troubleshoot your systems and applications using your existing system, application and custom log files. With CloudWatch Logs, you can monitor your logs, in near real time, for specific phrases, values or patterns.

Question 18. What Kinds Of Things Can I Do With My Logs In Cloudwatch?

Answer :

CloudWatch Logs is capable of monitoring and storing your logs to help you better understand and operate your systems and applications. When you use CloudWatch Logs with your logs, your existing log data is used for monitoring, so no code changes are required.

Question 19. What Types Of Data Can I Send To Cloudwatch Logs From My Ec2 Instances Running Sql Server And Windows Server?

Answer :

You can configure the EC2Config service to send a variety of data and log files to CloudWatch including: custom text logs, Even logs, Event Tracing logs, and Performance Counter data.

Question 20. How Frequently Does The Cloudwatch Logs Agent Send Data?

Answer :

The CloudWatch Logs Agent will send log data files every five seconds by default and is configurable by the user.

Question 21. What If I Configure The Cloudwatch Logs Agent To Send Non-text Log Data?

Answer :

The CloudWatch Logs Agent will record an error in the event it has been configured to report non-text log data.

Question 22. What Is The Syntax Of Metric Filter Patterns?

Answer :

A Metric Filter pattern can contain search terms or a specification of your common log or JSON event format. CloudWatch Logs can also be used to extract values from a log even in common log or JSON format.

Question 23. Can I Use Regular Expressions With My Log Data?

Answer :

CloudWatch Metric Filters does not support regular expressions. To process your log data with regular expressions, consider using Amazon Kinesis and connect the stream with a regular expression processing engine.

Question 24. How Do I Retrieve My Log Data?

Answer :

You can retrieve any of your log data using the CloudWatch Logs console or through the CloudWatch Logs CLI. Log events are retrieved based on the Log Group, Log Stream and time with which they are associated.

Question 25. How Long Does Cloudwatch Logs Store My Data?

Answer :

You can store your log data for as long as you want. By default, CloudWatch Logs will store your log data indefinitely. You can change the retention for each LogGroup at any time.

Question 26. What Types Of Cloudwatch Alarms Can Be Created?

Answer :

You can create an alarm to monitor any CloudWatch metric in your account. You can also create an alarm on custom metrics that are specific to your custom application or infrastructure.

Question 27. What Actions Can I Take From A Cloudwatch Alarm?

Answer :

When you create an alarm, you can configure it to perform one or more automated actions when the metric you chose to monitor exceeds a threshold you define.

Ex: send an email, publish to SQS, stop/terminate an EC2 instance, or execute an auto-scaling policy.

Question 28. What Thresholds Can I Set To Trigger A Cloudwatch Alarm?

Answer :

When you create an alarm, you first choose the CloudWatch metric you want it to monitor. Next, you choose the evaluation period and a statistical value to measure. To set a threshold, set a target value and choose whether the alarm will trigger when the value is greater, equal, or less than that value[8/6, 09:00] +91 96866 77992: Q1) What is AWS?

AWS stands for Amazon Web Services. AWS is a platform that provides on-demand resources

for hosting web services, storage, networking, databases and other resources over the internet with a pay-as-you-go pricing.

Q2) What are the components of AWS?

EC2 – Elastic Compute Cloud, S3 – Simple Storage Service, Route53, EBS – Elastic Block Store, Cloudwatch, Key-Pairs are few of the components of AWS.

Q3) What are key-pairs?

Key-pairs are secure login information for your instances/virtual machines. To connect to the instances we use key-pairs that contain a public-key and private-key.

Q4) What is S3?

S3 stands for Simple Storage Service. It is a storage service that provides an interface that you can use to store any amount of data, at any time, from anywhere in the world. With S3 you pay only for what you use and the payment model is pay-as-you-go.

Q5) What are the pricing models for EC2 instances?

The different pricing model for EC2 instances are as below,

On-demand Reserved Spot Scheduled Dedicated
Q6) What are the types of volumes for EC2 instances?
There are two types of volumes, Instance store volumes
EBS – Elastic Block Stores
Q7) What are EBS volumes?

EBS stands for Elastic Block Stores. They are persistent volumes that you can attach to the instances. With EBS volumes, your data will be preserved even when you stop your instances,

unlike your instance store volumes where the data is deleted when you stop the instances.

Q8) What are the types of volumes in EBS?

Following are the types of volumes in EBS,

General purpose Provisioned IOPS Magnetic Cold HDD Throughput optimized Q9) What are the different types of instances?

Following are the types of instances,

General purpose Computer Optimized Storage Optimized Memory Optimized Accelerated Computing Q10) What is an auto-scaling and what are the components?

Auto scaling allows you to automatically scale-up and scale-down the number of instances depending on the CPU utilization or memory utilization. There are 2 components in Auto scaling, they are Auto-scaling groups and Launch Configuration.

Q11) What are reserved instances?

Reserved instances are the instance that you can reserve a fixed capacity of EC2 instances. In reserved instances you will have to get into a contract of 1 year or 3 years.

Q12) What is an AMI?

AMI stands for Amazon Machine Image. AMI is a template that contains the software configurations, launch permission and a block device mapping that specifies the volume to attach to the instance when it is launched.

Q13) What is an EIP?

EIP stands for Elastic IP address. It is designed for dynamic cloud computing. When you want to have a static IP address for your instances when you stop and restart your instances, you will be using EIP address.

Q14) What is Cloudwatch?

Cloudwatch is a monitoring tool that you can use to monitor your various AWS resources. Like health check, network, Application, etc.

Q15) What are the types in cloudwatch?

There are 2 types in cloudwatch. Basic monitoring and detailed monitoring. Basic monitoring is free and detailed monitoring is chargeable.

Q16) What are the cloudwatch metrics that are available for EC2 instances?

Diskreads, Diskwrites, CPU utilization, networkpacketsIn, networkpacketsOut, networkIn, networkOut, CPUCreditUsage, CPUCreditBalance.

Q17) What is the minimum and maximum size of individual objects that you can store in S3

The minimum size of individual objects that you can store in S3 is 0 bytes and the maximum bytes that you can store for individual objects is 5TB.

Q18) What are the different storage classes in S3?

Following are the types of storage classes in S3,

Standard frequently accessed Standard infrequently accessed One-zone infrequently accessed Glacier RRS – reduced redundancy storage Q19) What is the default storage class in S3?

The default storage class in S3 is Standard frequently accessed.

Q20) What is glacier?

Glacier is the back up or archival tool that you use to back up your data in S3.

Q21) How can you secure the access to your S3 bucket?

There are two ways that you can control the access to your S3 buckets,

ACL – Access Control List Bucket policies Q22) How can you encrypt data in S3?

You can encrypt the data by using the below methods,

Server Side Encryption – S3 (AES 256 encryption) Server Side Encryption – KMS (Key management Service) Server Side Encryption – C (Client Side) Q23) What are the parameters for S3 pricing?

The pricing model for S3 is as below,

Storage used
Number of requests you make
Storage management
Data transfer
Transfer acceleration

Q24) What is the pre-requisite to work with Cross region replication in S3?

You need to enable versioning on both source bucket and destination to work with cross region replication. Also both the source and destination bucket should be in different region.

Q25) What are roles?

Roles are used to provide permissions to entities that you trust within your AWS account. Roles are users in another account. Roles are similar to users but with roles you do not need to create any username and password to work with the resources.

Q26) What are policies and what are the types of policies?

Policies are permissions that you can attach to the users that you create. These policies will contain that access that you have provided to the users that you have created. There are 2 types of policies.

Managed policies
Inline policies

Q27) What is cloudfront?

Cloudfront is an AWS web service that provided businesses and application developers an easy and efficient way to distribute their content with low latency and high data transfer speeds. Cloudfront is content delivery network of AWS.

Q28) What are edge locations?

Edge location is the place where the contents will be cached. When a user tries to access some content, the content will be searched in the edge location. If it is not available then the content will be made available from the origin location and a copy will be stored in the edge location.

Q29) What is the maximum individual archive that you can store in glacier?

You can store a maximum individual archive of upto 40 TB.

Q30) What is VPC?

VPC stands for Virtual Private Cloud. VPC allows you to easily customize your networking configuration. VPC is a network that is logically isolated from other network in the cloud. It allows you to have your own IP address range, subnets, internet gateways, NAT gateways and security groups.

Q31) What is VPC peering connection?

VPC peering connection allows you to connect 1 VPC with another VPC. Instances in these VPC behave as if they are in the same network.

Q32) What are NAT gateways?

NAT stands for Network Address Translation. NAT gateways enables instances in a private subnet to connect to the internet but prevent the internet from initiating a connection with those instances.

Q33) How can you control the security to your VPC?

You can use security groups and NACL (Network Access Control List) to control the security to your

VPC.

Q34) What are the different types of storage gateway?

Following are the types of storage gateway.

File gateway Volume gateway Tape gateway Q35) What is a snowball?

Snowball is a data transport solution that used source appliances to transfer large amounts of data into and out of AWS. Using snowball, you can move huge amount of data from one place to another which reduces your network costs, long transfer times and also provides better security.

Q36) What are the database types in RDS?

Following are the types of databases in RDS,

Aurora Oracle MySQL server PostgreSQL MariaDB SQL server Q37) What is a redshift?

Amazon redshift is a data warehouse product. It is a fast and powerful, fully managed, petabyte scale data warehouse service in the cloud.

Q38) What is SNS?

SNS stands for Simple Notification Service. SNS is a web service that makes it easy to notifications from the cloud. You can set up SNS to receive email notification or message notification.

Q39) What are the types of routing policies in route53?

Following are the types of routing policies in route53,

Simple routing Latency routing Failover routing Geolocation routing Weighted routing Multivalue
answer Q40) What is the maximum size of messages in SQS?

The maximum size of messages in SQS is 256 KB.

Q41) What are the types of queues in SQS?

There are 2 types of queues in SQS.

Standard queue FIFO (First In First Out) Q42) What is multi-AZ RDS?

Multi-AZ (Availability Zone) RDS allows you to have a replica of your production database in another availability zone. Multi-AZ (Availability Zone) database is used for disaster recovery. You will have an exact copy of your database. So when your primary database goes down, your application will automatically failover to the standby database.

Q43) What are the types of backups in RDS database?

There are 2 types of backups in RDS database.

Automated backups Manual backups which are known as snapshots. Q44) What is the difference between security groups and network access control list? Security Groups Network access control list Can control the access at the instance level Can control access at the subnet level Can add rules for "allow" only Can add rules for both "allow" and "deny" Evaluates all rules before allowing the traffic Rules are processed in order number when allowing traffic. Can assign

unlimited number of security groups Can assign upto 5 security groups. Stateful filtering Stateless filtering Q45) What are the types of load balancers in EC2?

There are 3 types of load balancers,

Application load balancer Network load balancer Classic load balancer Q46) What is an ELB?

ELB stands for Elastic Load balancing. ELB automatically distributes the incoming application traffic or network traffic across multiple targets like EC2, containers, IP addresses.

Q47) What are the two types of access that you can provide when you are creating users?

Following are the two types of access that you can create.

Programmatic access Console access Q48) What are the benefits of auto scaling?

Following are the benefits of auto scaling

Better fault tolerance Better availability Better cost management Q49) What are security groups?

Security groups act as a firewall that contains the traffic for one or more instances. You can associate one or more security groups to your instances when you launch them. You can add rules to each security group that allow traffic to and from its associated instances. You can modify the rules of a security group at any time, the new rules are automatically and immediately applied to all the instances that are associated with the security group

Q50) What are shared AMIs?

Shared AMI's are the AMI that are created by other developed and made available for other developed to use.

Q51) What is the difference between the classic load balancer and application load balancer?

Answer: Dynamic port mapping, multiple port multiple listeners is used in Application Load Balancer, One port one listener is achieved via Classic Load Balancer

Q52) By default how many Ip address does aws reserve in a subnet?

Answer: 5

Q53) What is meant by subnet?

Answer: A large section of IP Address divided in to chunks are known as subnets

Q54) How can you convert a public subnet to private subnet?

Answer: Remove IGW & add NAT Gateway, Associate subnet in Private route table

Q55) Is it possible to reduce a ebs volume?

Answer: no it's not possible, we can increase it but not reduce them

Q56) What is the use of elastic ip are they charged by AWS?

Answer: These are ipv4 address which are used to connect the instance from internet, they are charged if the instances are not attached to it

Q57) One of my s3 bucket is deleted but i need to restore is there any possible way?

Answer: If versioning is enabled we can easily restore them

Q58) When I try to launch an ec2 instance i am getting Service limit exceed, how to fix the issue?

Answer: By default AWS offer service limit of 20 running instances per region, to fix the issue we need to contact AWS support to increase the limit based on the requirement

Q59) I need to modify the ebs volumes in Linux and windows is it possible

Answer: yes its possible from console use modify volumes in section give the size u need then for windows go to disk management for Linux mount it to achieve the modification

Q60) Is it possible to stop a RDS instance, how can I do that?

Answer: Yes it's possible to stop rds. Instance which are non-production and non multi AZ's

Q61) What is meant by parameter groups in rds. And what is the use of it?

Answer: Since RDS is a managed service AWS offers a wide set of parameter in RDS as parameter group which is modified as per requirement

Q62) What is the use of tags and how they are useful?

Answer: Tags are used for identification and grouping AWS Resources

Q63) I am viewing an AWS Console but unable to launch the instance, I receive an IAM Error how can I rectify it?

Answer: As AWS user I don't have access to use it, I need to have permissions to use it further

Q64) I don't want my AWS Account id to be exposed to users how can I avoid it?

Answer: In IAM console there is option as sign in url where I can rename my own account name with AWS account

Q65) By default how many Elastic Ip address does AWS Offer?

Answer: 5 elastic ip per region

Q66) You are enabled sticky session with ELB. What does it do with your instance?

Answer: Binds the user session with a specific instance

Q67) Which type of load balancer makes routing decisions at either the transport layer or the Application layer and supports either EC2 or VPC.

Answer: Classic Load Balancer

Q68) Which is virtual network interface that you can attach to an instance in a VPC?

Answer: Elastic Network Interface

Q69) You have launched a Linux instance in AWS EC2. While configuring security group, you have selected SSH, HTTP, HTTPS protocol. Why do we need to select SSH?

Answer: To verify that there is a rule that allows traffic from EC2 Instance to your computer

Q70) You have chosen a windows instance with Classic and you want to make some change to the security group. How will these changes be effective?

Answer: Changes are automatically applied to windows instances

Q71) Load Balancer and DNS service comes under which type of cloud service?

Answer: IaaS-Storage

Q72) You have an EC2 instance that has an unencrypted volume. You want to create another Encrypted volume from this unencrypted volume. Which of the following steps can achieve this?

Answer: Create a snapshot of the unencrypted volume (applying encryption parameters), copy the snapshot and create a volume from the copied snapshot

Q73) Where does the user specify the maximum number of instances with the auto scaling Commands?

Answer: Auto scaling Launch Config

Q74) Which are the types of AMI provided by AWS?

Answer: Instance Store backed, EBS Backed

Q75) After configuring ELB, you need to ensure that the user requests are always attached to a Single instance. What setting can you use?

Answer: Sticky session

Q76) When do I prefer to Provisioned IOPS over the Standard RDS storage?

Ans: If you have do batch-oriented is workloads.

Q77) If I am running on my DB Instance a Multi-AZ deployments, can I use to the stand by the DB Instance for read or write a operation along with to primary DB instance?

Ans: Primary db instance does not working.

Q78) Which the AWS services will you use to the collect and the process e-commerce data for the near by real-time analysis?

Ans: Good of Amazon DynamoDB.

Q79) A company is deploying the new two-tier an web application in AWS. The company has to limited on staff and the requires high availability, and the application requires to complex queries and table joins. Which configuration provides to the solution for company's requirements?

Ans: An web application provide on Amazon DynamoDB solution.

Q80) Which the statement use to cases are suitable for Amazon DynamoDB?

Ans: The storing metadata for the Amazon S3 objects& The Running of relational joins and complex an updates.

Q81) Your application has to the retrieve on data from your user's mobile take every 5 minutes and then data is stored in the DynamoDB, later every day at the particular time the data is an extracted into S3 on a per user basis and then your application is later on used to visualize the data to user. You are the asked to the optimize the architecture of the backend system can to lower cost, what would you recommend do?

Ans: Introduce Amazon ElastiCache to the cache reads from the Amazon DynamoDB table and to reduce the provisioned read throughput.

Q82) You are running to website on EC2 instances can deployed across multiple Availability Zones with an Multi-AZ RDS MySQL Extra Large DB Instance etc. Then site performs a high number of the small reads and the write per second and the relies on the eventual consistency model. After the comprehensive tests you discover to that there is read contention on RDS MySQL. Which is the best approaches to the meet these requirements?

Ans:The Deploy Elasti Cache in-memory cache is running in each availability zone and Then Increase the RDS MySQL Instance size and the Implement provisioned IOPS.

Q83) An startup is running to a pilot deployment of around 100 sensors to the measure street noise and The air quality is urban areas for the 3 months. It was noted that every month to around the 4GB of sensor data are generated. The company uses to a load balanced take auto scaled layer of the EC2 instances and a RDS database with a 500 GB standard storage. The pilot was success and now they want to the deploy take atleast 100K sensors.let which to need the supported by backend. You need to the stored data for at least 2 years to an analyze it. Which

setup of following would you be prefer?

Ans: The Replace the RDS instance with an 6 node Redshift cluster with take 96TB of storage.

Q84) Let to Suppose you have an application where do you have to render images and also do some of general computing. which service will be best fit your need?

Ans: Used on Application Load Balancer.

Q85) How will change the instance give type for the instances, which are the running in your applications tier and Then using Auto Scaling. Where will you change it from areas?

Ans: Changed to Auto Scaling launch configuration areas.

Q86) You have an content management system running on the Amazon EC2 instance that is the approaching 100% CPU of utilization. Which option will be reduce load on the Amazon EC2 instance?

Ans: Let Create a load balancer, and Give register the Amazon EC2 instance with it.

Q87) What does the Connection of draining do?

Ans: The re-routes traffic from the instances which are to be updated (or) failed an health to check.

Q88) When the instance is an unhealthy, it is do terminated and replaced with an new ones, which of the services does that?

Ans: The service makes a fault tolerance.

Q89) What are the lifecycle hooks used for the AutoScaling?

Ans: They are used to put an additional taken wait time to the scale in or scale out events.

Q90) An user has to setup an Auto Scaling group. Due to some issue the group has failed to launch a single instance for more than 24 hours. What will happen to the Auto Scaling in the condition?

Ans: The auto Scaling will be suspended to the scaling process.

Q91) You have an EC2 Security Group with several running EC2 instances. You changed the Security Group rules to allow the inbound traffic on a new port and protocol, and then launched several new instances in the same Security Group. Such the new rules apply?

Ans: The immediately to all the instances in security groups.

[8/6, 09:00] +91 96866 77992: What is Amazon Cloudwatch?

Answer :

CloudWatch is a monitoring service for AWS cloud resources and the applications you run on AWS. You can use CloudWatch to collect and track metrics, collect and monitor log files, and set alarms. CloudWatch can monitor resources such as EC2 instances, DynamoDB tables, and RDS DB instances.

Question 2. Which Operating Systems Does Cloudwatch Support?

Answer :

CloudWatch receives and provides metrics for all EC2 instances and should work with any operating system currently supported by the EC2 service.

Question 3. What Access Management Policies Can I Implement For Cloudwatch?

Answer :

CloudWatch integrates with AWS IAM so that you can specify which CloudWatch actions a user in your AWS Account can perform. You can't use IAM to control access to CloudWatch data for specific resources. For example, you can't give a user access to CloudWatch data for only a specific set of instances or a specific LoadBalancer. Permissions granted using IAM cover all the cloud resources you use with CloudWatch. In addition, you can't use IAM roles with the Amazon CloudWatch command line tools.

Question 4. What Is Amazon Cloudwatch Logs?

Answer :

Amazon CloudWatch Logs lets you monitor and troubleshoot your systems and applications using your existing system, application and custom log files. With CloudWatch Logs, you can monitor your logs, in near real time, for specific phrases, values or patterns. For example, you could set an alarm on the number of errors that occur in your system logs or view graphs of latency of web requests from your application logs. You can then view the original log data to see the source of the problem. Log data can be stored and accessed indefinitely in highly durable, low-cost storage so you don't have to worry about filling up hard drives.

Question 5. What Kinds Of Things Can I Do With Cloudwatch Logs?

Answer :

CloudWatch Logs is capable of monitoring and storing your logs to help you better understand and operate your systems and applications.

You can use CloudWatch logs in a number of ways:

Real Time application and system monitoring, Long term log retention.

[8/6, 09:00] +91 96866 77992: uestion 6. What Platforms Does The Cloudwatch Logs Agent Support?

Answer :

The CloudWatch Log Agent is supported on Amazon Linux, Ubuntu, CentOS, Red Hat Enterprise Linux, and Windows. This agent will support the ability to monitor individual log files on the host.

Question 7. Does The Cloudwatch Logs Agent Support Iam Roles?

Answer :

Yes. The CloudWatch Logs Agent is integrated with IAM and includes support for both access keys and IAM roles.

Question 8. Does The Amazon Cloudwatch Monitoring Charge Change Depending On Which Ec2 Instance I Monitor?

Answer :

No, CloudWatch monitoring charge does not vary by EC2 instance type.

Question 9. What Can I Measure With Amazon Cloudwatch Metrics?

Answer :

CloudWatch allows you to monitor AWS cloud resources and the applications you run on AWS. Metrics are provided automatically for a number of AWS products and services, including EC2 instances, EBS volumes, ELBs, Autoscaling groups, EMR job flows, RDS DB instances, DynamoDB tables, ElastiCache clusters, RedShift clusters, OpsWorks stacks, Route 53 health checks, SNS topics, SQS queues, SWF workflows, and Storage Gateways. You can also monitor custom metrics generated by your own applications and services.

Question 10. What Is The Retention Period Of All Metrics?

Answer :

CloudWatch retains metric data as follows:

Data points with a period of less than 60 seconds are available for 3 hours. These data points are high-resolution custom metrics, Data points with a period of 60 seconds (1 minute) are available for 15 days, Data points with a period of 300 seconds (5 minute) are available for 63 days, 4) Data points with a metric of 3600 seconds (1 hr) are available for 455 days (15 months). Data points that are initially published with a shorter period are aggregated together for long-term storage.

[8/6, 09:01] +91 96866 77992: Will I Lose The Metrics Data If I Disable Monitoring For An Ec2 Instance?

Answer :

No. You can always retrieve metrics data for any EC2 instance based on the retention schedules. However, the CloudWatch console limits the search of metrics to 2 weeks after a metric is last ingested to ensure that the most up to date instances are shown in your namespace.

Question 12. Can I Access The Metrics Data For A Terminated Ec2 Instance Or A Deleted Elb?

Answer :

Yes, CloudWatch stores metrics for terminated EC2 instances or deleted ELBs for 15 months.

Question 13. What Metrics Are Available At High Resolution?

Answer :

Currently, only customer metrics that you publish to CloudWatch are available at high resolution.

Question 14. Are High-resolution Custom Metrics Priced Differently Than Regular Custom Metrics?

Answer :

No, high-resolution custom metrics are priced in the same manner as standard 1-min custom metrics.

Question 15. When Would I Use A Custom Metric Over Having My Program Emit A Log To Cloudwatch Logs?

Answer :

You can monitor your own data using custom metrics, CloudWatch logs, or both. You may want to use custom metrics if your data is not already produced in log format, for example OS process or performance measurements. Or, you may want to write your own application or script, or one provided by an AWS partner. If you want to store and save individual measurements along with additional detail, you may want to use CloudWatch Logs.

Question 16. What Statistics Can I View And Graph In Cloudwatch?

Answer :

You can retrieve, graph, and set alarms on the following statistical metrics: average, sum, min, max, and sample count.

[8/6, 09:01] +91 96866 77992: 17. What Log Monitoring Does Cloudwatch Provide?

Answer :

CloudWatch Logs lets you monitor and troubleshoot your systems and applications using your existing system, application and custom log files. With CloudWatch Logs, you can monitor your logs, in near real time, for specific phrases, values or patterns.

Question 18. What Kinds Of Things Can I Do With My Logs In Cloudwatch?

Answer :

CloudWatch Logs is capable of monitoring and storing your logs to help you better understand and operate your systems and applications. When you use CloudWatch Logs with your logs,

your existing log data is used for monitoring, so no code changes are required.

Question 19. What Types Of Data Can I Send To Cloudwatch Logs From My Ec2 Instances Running Sql Server And Windows Server?

Answer :

You can configure the EC2Config service to send a variety of data and log files to CloudWatch including: custom text logs, Even logs, Event Tracing logs, and Performance Counter data.

Question 20. How Frequently Does The Cloudwatch Logs Agent Send Data?

Answer :

The CloudWatch Logs Agent will send log data files every five seconds by default and is configurable by the user.

Question 21. What If I Configure The Cloudwatch Logs Agent To Send Non-text Log Data?

Answer :

The CloudWatch Logs Agent will record an error in the event it has been configured to report non-text log data.

Question 22. What Is The Syntax Of Metric Filter Patterns?

Answer :

A Metric Filter pattern can contain search terms or a specification of your common log or JSON event format. CloudWatch Logs can also be used to extract values from a log even in common log or JSON format.

Question 23. Can I Use Regular Expressions With My Log Data?

Answer :

CloudWatch Metric Filters does not support regular expressions. To process your log data with regular expressions, consider using Amazon Kinesis and connect the stream with a regular expression processing engine.

Question 24. How Do I Retrieve My Log Data?

Answer :

You can retrieve any of your log data using the CloudWatch Logs console or through the CloudWatch Logs CLI. Log events are retrieved based on the Log Group, Log Stream and time with which they are associated.

Question 25. How Long Does Cloudwatch Logs Store My Data?

Answer :

You can store your log data for as long as you want. By default, CloudWatch Logs will store your log data indefinitely. You can change the retention for each LogGroup at any time.

Question 26. What Types Of Cloudwatch Alarms Can Be Created?

Answer :

You can create an alarm to monitor any CloudWatch metric in your account. You can also create an alarm on custom metrics that are specific to your custom application or infrastructure.

Question 27. What Actions Can I Take From A Cloudwatch Alarm?

Answer :

When you create an alarm, you can configure it to perform one or more automated actions when the metric you chose to monitor exceeds a threshold you define.

Ex: send an email, publish to SQS, stop/terminate an EC2 instance, or execute an auto-scaling policy.

Question 28. What Thresholds Can I Set To Trigger A Cloudwatch Alarm?

Answer :

When you create an alarm, you first choose the CloudWatch metric you want it to monitor. Next, you choose the evaluation period and a statistical value to measure. To set a threshold, set a target value and choose whether the alarm will trigger when the value is greater, equal, or less than that value